

# NNS 系列之 Registry

---

## Contents

---

- [NNS 系列之 Registry](#)
- [Contents](#)
  - [概述](#)
  - [数据结构](#)
    - [REGISTRY](#)
  - [接口函数](#)
    - [canister\\_init](#)
    - [canister\\_pre\\_upgrade](#)
    - [canister\\_post\\_upgrade](#)
    - [update\\_authz](#)
    - [current\\_authz](#)
    - [get\\_changes\\_since](#)
    - [get\\_certified\\_changes\\_since](#)
    - [get\\_value](#)
    - [get\\_latest\\_version](#)
    - [get\\_certified\\_latest\\_version](#)
    - [atomic\\_mutate](#)
    - [bless\\_replica\\_version](#)
    - [update\\_subnet\\_replica\\_version](#)
    - [update\\_icp\\_xdr\\_conversion\\_rate](#)
    - [add\\_node](#)
    - [add\\_node\\_operator](#)
    - [create\\_subnet](#)
    - [add\\_nodes\\_to\\_subnet](#)
    - [delete\\_subnet](#)
    - [recover\\_subnet](#)
    - [remove\\_nodes\\_from\\_subnet](#)
    - [remove\\_nodes](#)
    - [remove\\_node\\_directly](#)
    - [update\\_node\\_operator\\_config](#)
    - [update\\_subnet](#)
    - [clear\\_provisional\\_whitelist](#)
    - [set\\_firewall\\_config](#)
    - [update\\_node\\_rewards\\_table](#)
    - [add\\_or\\_remove\\_data\\_centers](#)
    - [update\\_unassigned\\_nodes\\_config](#)
    - [get\\_node\\_providers\\_monthly\\_xdr\\_rewards](#)
- [Reference](#)

## 概述

Registry canister，功能上是执行。在 Governance 得到投票结果之后，Governance 就会调用 Registry 里面的 Functions 来具体执行。

## 数据结构

### REGISTRY

Registry					
version:	Version(u64)				
pub store:	RegistryMap/BTreeMap				
	K=	Vec <u8< u=""></u8<>	RegistryValue		
	V=	VecDeque	pub value: Vec <u8< u=""></u8<>		
			pub version: u64		
pub deletion_marker: bool					
changelog:	RbTree<K, V>				
	root:	NodeRef<K, V>			
		Option<Box	Node<K, V>		
			key:	K=	EncodedVersion
					[u8; 8]
				value:	V= Vec <u8< u=""></u8<>
			left:	NodeRef<K, V>	
			right:	NodeRef<K, V>	
			color:	Color(enum)	
				Red	
Black					
subtree	Hash				
	_hash:	[u8; 32]			

- 核心的数据结构就是 Registry。它有三个字段：
  - version，一个全局的计数器，每次有 mutation 作用于 Registry，就会 +1。
  - store 是真正存数据的地方，一个 Map，它的 Key 和 Value 需要转化成 u8 (bytes) 数组。
  - changelog 所有的 mutation 的历史记录。

接下来看一下怎么示例化的。

实现了一个全局唯一的静态变量，**REGISTRY**，来存储所有的状态。

并且通过两个函数，`registry()` 以及 `registry_mut()` 来获取到这个变量。它们的区别是上面的 `registry()` 获取到的 `REGISTRY` 只读，不能更改它的状态，下面这个 `registry_mut()` 还能更改其状态。

```
static mut REGISTRY: Option<Registry> = None;

fn registry() -> &'static Registry {
    registry_mut()
}

fn registry_mut() -> &'static mut Registry {
    unsafe {
        if let Some(g) = &mut REGISTRY {
            g
        } else {
            REGISTRY = Some(Registry::new());
            registry_mut()
        }
    }
}
```

## 接口函数

### canister\_init

在 `canister` 部署的时候，需要初始化一些状态，执行初始化函数，和 Solidity 的 constructor（构造函数）类似。比如部署一个 token 合约，需要初始化 token 的名字，符号，总量等信息。这就是 `canister_init` 函数的功能。

先执行了 `hook()` 函数，`hook()` 来自于 `dfn_core` 这个包里面，去看一下它的具体内容，它执行了两个函数，`set_stdout()` 和 `set_panic_hook()`。根据 `set_panic_hook()` 的具体内容，总的来说，`hook()` 它的功能是设置好标准输出，以及程序出现 Panic 时，把错误消息，以及所在文件，第几行，第几列的 log 打印出来。

```
//
https://github.com/dfinity/ic/blob/8fffeb4be1/rs/registry/canister/canister/canister.rs#L118
dfn_core::printer::hook();

//
https://github.com/dfinity/ic/blob/8fffeb4be1/rs/rust\_canisters/dfn\_core/src/printer.rs
/// Sets stdout, stderr, and a custom panic hook
pub fn hook() {
    set_stdout();
    set_panic_hook();
}

//
```

```

https://github.com/dfinity/ic/blob/8fffeb4be1/rs/rust_canisters/dfn_core/s
rc/printer.rs#L118-L146
/// Sets a custom panic hook, uses debug.trace
pub fn set_panic_hook() {
    panic::set_hook(Box::new(|info| {
        let file = info.location().unwrap().file();
        let line = info.location().unwrap().line();
        let col = info.location().unwrap().column();

        let msg = match info.payload().downcast_ref:::<'static str>() {
            Some(s) => *s,
            None => match info.payload().downcast_ref:::<String>() {
                Some(s) => &s[..],
                None => "Box<Any>",
            },
        };

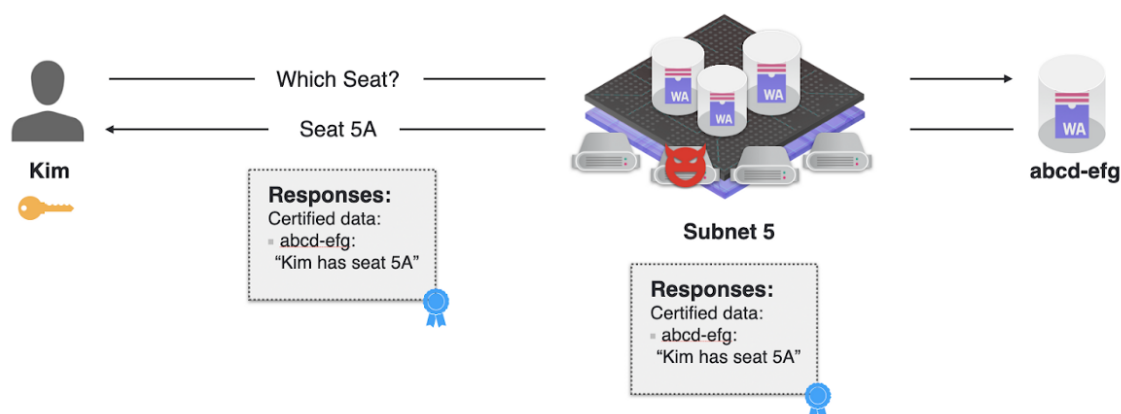
        let err_info = format!("Panicked at '{}', {}:{}:{}", msg, file,
line, col);

        unsafe {
            log(&err_info);
        }
        crate::api::trap_with(&err_info);
    }));
}

```

接下来执行 `recertify_registry()` 这个函数。先整体说一下 `certified_data`。

## Using Certified Variables: The Query Call



```
recertify_registry();
```

```
fn recertify_registry() {
    use ic_certified_map::{fork_hash, labeled_hash};

    let root_hash = fork_hash(
        &current_version_tree(registry().latest_version()).reconstruct(),
        &labeled_hash(b"delta", &registry().changelog().root_hash()),
    );
    set_certified_data(&root_hash);
}
```

先从里面的 `registry()` 开始看，它返回了一个只读的那个全局的 Registry 实例。然后去拿它的 `latest_version()`，其实就是拿的它 version 字段的值，当前版本号，一个 u64 的值。

```
//
https://github.com/dfinity/ic/blob/8ffffeb4be1/rs/registry/canister/src/registry.rs#L177-L179
pub fn latest_version(&self) -> Version {
    self.version
}
```

然后把这个 Version 输入到 `current_version_tree()` 里面，看一下 `current_version_tree()`，它接受一个 Version 类型，我们之前在 struct 看过，其实就是 u64，然后返回一个 HashTree 里面的 Labeled 节点，把 current\_version 存为数据标签，连接着一个 Leaf 叶子节点，里面存了传进去的 Version 的编码。(leb128 是 Little Endian Base 128 是编码算法，对于存储大整数有空间优势)

```
use ic_certified_map::{labeled, HashTree};

/// The maximum amount of bytes a 64-bit number can occupy when encoded in
/// LEB128.
const MAX_U64_ENCODING_BYTES: usize = 10;

pub fn current_version_tree(v: Version) -> HashTree<'static> {
    let mut buf = Vec::with_capacity(MAX_U64_ENCODING_BYTES);
    leb128::write::unsigned(&mut buf, v).unwrap();
    labeled(
        b"current_version",
        HashTree::Leaf(std::borrow::Cow::from(buf)),
    )
}
```

接下来需要看一下 HashTree 是怎么回事，以及包括 ic-certified-map crate（包）的函数：fork\_hash, labeled\_hash。

下面定义了一个 Hash, Sha256 的输出，32 位的 u8 数组。还定义了一个 HashTree 的数据结构，一个树结构，里面有几类节点，比如空节点；下面有两个子节点，它本身没有任何值的 Fork 节点；下面有一个节点，并且本身带标签（数据）的 Labeled 节点；或者只有数据的 Leaf 节点，或者是 [u8; 32] 的 Pruned 节点。

```
// https://github.com/dfinity/cdk-rs/blob/319795e9b4/src/ic-certified-
map/src/hashtree.rs#L9-L49

/// SHA-256 hash bytes.
pub type Hash = [u8; 32];

/// HashTree as defined in the interfaces spec.
/// https://sdk.dfinity.org/docs/interface-spec/index.html#_certificate
#[derive(Debug)]
pub enum HashTree<'a> {
    Empty,
    Fork(Box<(HashTree<'a>, HashTree<'a>)>),
    Labeled(&'a [u8], Box<HashTree<'a>>),
    Leaf(Cow<'a, [u8]>),
    Pruned(Hash),
}

pub fn fork<'a>(l: HashTree<'a>, r: HashTree<'a>) -> HashTree<'a> {
    HashTree::Fork(Box::new((l, r)))
}

pub fn labeled<'a>(l: &'a [u8], t: HashTree<'a>) -> HashTree<'a> {
    HashTree::Labeled(l, Box::new(t))
}

pub fn fork_hash(l: &Hash, r: &Hash) -> Hash {
    let mut h = domain_sep("ic-hashtree-fork");
    h.update(&l[..]);
    h.update(&r[..]);
    h.finalize().into()
}

pub fn leaf_hash(data: &[u8]) -> Hash {
    let mut h = domain_sep("ic-hashtree-leaf");
    h.update(data);
    h.finalize().into()
}

pub fn labeled_hash(label: &[u8], content_hash: &Hash) -> Hash {
    let mut h = domain_sep("ic-hashtree-labeled");
    h.update(label);
    h.update(&content_hash[..]);
    h.finalize().into()
}

fn domain_sep(s: &str) -> sha2::Sha256 {
    let buf: [u8; 1] = [s.len() as u8];
    let mut h = Sha256::new();
    h.update(&buf[..]);
    h.update(s.as_bytes());
    h
}

```

所以 `fork_hash`, `labeled_hash` 都是输入不可变量, 通过标准的 Sha256, 然后往 Hash 的 buf 里面, 填充一些自定义字段, 比如 `ic-hashtree-fork` 这样的字符串。最后输出 32 位的 bytes 数组。

HashTree 实现了一个方法, `reconstruct()` 其中输入一个 HashTree, 输出一个 Hash。Hash 就是 32 位的 Bytes 数组。`reconstruct()` 这个函数, 和前面的 `fork_hash()`, `labeled_hash()`, `leaf_hash()` 相呼应, 总的来说, 将一个 HashTree 这样的树, 递归的 hash 了一遍, 并最后输出一个 32 位的 u8 数组。

```
impl HashTree<'_> {
    pub fn reconstruct(&self) -> Hash {
        match self {
            Self::Empty => domain_sep("ic-hashtree-
empty").finalize().into(),
            Self::Fork(f) => fork_hash(&f.0.reconstruct(),
&f.1.reconstruct()),
            Self::Labeled(l, t) => {
                let thash = t.reconstruct();
                labeled_hash(l, &thash)
            }
            Self::Leaf(data) => leaf_hash(data),
            Self::Pruned(h) => *h,
        }
    }
}
```

可以再去我们之前的 Struct 图看一下, RbTree 是怎么定义的。这里给它实现了一个 `root_hash()` 的方法, 其实就是拿的它 `subtree_hash`。

```
// https://github.com/dfinity/cdk-rs/blob/319795e9b4/src/ic-certified-
map/src/hashtree.rs#L57-L63
impl<K: 'static + AsRef<[u8]>, V: AsHashTree + 'static> AsHashTree for
RbTree<K, V> {
    fn root_hash(&self) -> Hash {
        match self.root.as_ref() {
            None => Empty.reconstruct(),
            Some(n) => n.subtree_hash,
        }
    }
}
```

然后 `labeled_hash` 输入一个 "delta" 的字符串的 bytes 形式, 以及 changelog 的 `subtree_hash`, 将结果和前面的到的 `reconstruct()` 的结果输入到 `fork_hash`, 最后得到 `root_hash`。

```
set_certified_data(&root_hash);

//
https://github.com/dfinity/ic/blob/8fffeb4be1/rs/rust_canisters/dfn_core/s
rc/api.rs#L840-L847
/// Sets the certified data of this canister.
```

```

///
/// # Panics
///
/// * This function traps if data.len() > 32.
pub fn set_certified_data(data: &[u8]) {
    // https://smartcontracts.org/docs/interface-spec/index.html#system-api-certified-data
    unsafe { ic0::certified_data_set(data.as_ptr() as u32, data.len() as u32) }
}

```

然后调用系统接口 `certified_data_set` 把这个 `root_hash` 设置成认证数据。认证数据的具体内容可以看 `interface-spec` 文档，也可以看[官方视频](#)，简单的说是你 query 一个数据的时候，没有经过全网共识，直接从一个节点拿的，可能该节点作恶。`certified_data` 是之前通过共识过放进去的，下次被 query 时还能携带一个全网的签名，就可以验证这个数据的安全性很高的。好像 `Canister Signature` 也和这个东西有关，期待有人更深入的分享 `certified_data` 相关。

所以总的来说，`recertify_registry()` 只用了当前 `Registry` 的 `version`，`changelog` 这两个字段，经过一系列的 hash，生成了 32 位的 Hash 结果，`root_hash`。`root_hash` 是能够反映 `Registry` 在 `version` 和 `changelog` 上的变化的。最后，这个 `root_hash` 设置成 `certified_data`。`Registry` 里面的 `store` 字段并没有影响到 `root_hash`，后续可以看出每次 `store` 变化，都会往 `changelog` 里面存这次变化，所以可能 `changelog` 里面的数据已经能够反映 `store` 的信息了。

接下来：

```

let init_payload =
    Decode!(&arg_data(), RegistryCanisterInitPayload)
    .expect("The init argument for the registry canister must be a Candid-encoded RegistryCanisterInitPayload.");
println!(
    "{}canister_init: Initializing with: {}",
    LOG_PREFIX, init_payload
);

//
// https://github.com/dfinity/ic/blob/8fffeb4be1/rs/rust_canisters/dfn_core/src/api.rs#L656-L664
// Returns the argument extracted from the message payload.
pub fn arg_data() -> Vec<u8> {
    let len: u32 = unsafe { ic0::msg_arg_data_size() };
    let mut bytes = vec![0; len as usize];
    unsafe {
        ic0::msg_arg_data_copy(bytes.as_mut_ptr() as u32, 0, len);
    }
    bytes
}

```



然后通过几个系统接口，`msg_arg_data_size` 和 `msg_arg_data_copy`，获取到这个函数的输入参数，因为是 `canister_init`，所以也是这个 `canister` 的安装参数。这个参数是 `RegistryCanisterInitPayload` 类型。

接着申明并初始化 `init_payload` 变量为 `Vec<u8>` 形式的参数解码成 `RegistryCanisterInitPayload` 类型的数据。

<b><i>RegistryCanisterInitPayload</i></b>		
<b><i>mutations: Vec&lt;</i></b>	<b><i>RegistryAtomicMutateRequest</i></b>	
	<b><i>mutations: Vec&lt;</i></b>	<b><i>RegistryMutation</i></b>
		<b><i>mutation_type: i32</i></b>
		<b><i>key: Vec&lt;u8&gt;</i></b>
	<b><i>preconditions: Vec&lt;</i></b>	<b><i>value: Vec&lt;u8&gt;</i></b>
		<b><i>Precondition</i></b>
		<b><i>key: Vec&lt;u8&gt;</i></b>
		<b><i>expected_version: Vec&lt;u8&gt;</i></b>

这个类型，首先有个唯一的字段 `mutations`，它是一个动态数组，而里面的 `RegistryAtomicMutateRequest` 有两个字段，其中一个也叫 `mutations`，另一个叫 `preconditions`。它们里面的字段都包含了一些 `Vec<u8>` 的值。

```
let registry = registry_mut();

init_payload
    .mutations
    .into_iter()
    .for_each(|mutation_request| {

registry.maybe_apply_mutation_internal(mutation_request.mutations)
    });
```

接着获取到了 `Registry` 示例，并且是可变的。接下来修改这个示例。

通过遍历 `init_payload` 的 `mutations` 这个数组里面的每一个 `RegistryAtomicMutateRequest` 类型的值，执行 `maybe_apply_mutation_internal` 方法，输入 `RegistryAtomicMutateRequest` 的 `mutations` 字段的数据 (`Vec<RegistryMutation>`) 来修改 `Registry` 示例。

这个方法执行了哪些功能？

```
//
https://github.com/dfinity/ic/blob/8ffffeb4be1/rs/registry/canister/src/registry.rs#L271-L294
/// Checks that invariants hold after applying mutations
pub fn maybe_apply_mutation_internal(&mut self, mutations:
Vec<RegistryMutation>) {
```

```
println!(
    "{}Received a mutate call containing a list of {} mutations",
    LOG_PREFIX,
    mutations.len()
);

let errors = self.verify_mutation_type(mutations.as_slice());
if !errors.is_empty() {
    panic!(
        "{}Transaction rejected because of the following errors:
[{}].",
        LOG_PREFIX,
        errors
            .iter()
            .map(|e| format!("{}", e))
            .collect::<Vec::<String>>()
            .join(", ")
    );
}

self.check_global_invariants(mutations.as_slice());
self.apply_mutations(mutations);
}
```

在做完一些验证和检查之后，真正的状态修改发生在 `apply_mutations()`

```
//
https://github.com/dfinity/ic/blob/8ffffeb4be1/rs/registry/canister/src/registry.rs#L226-L240
/// Applies the given mutations, without any check corresponding
/// to the mutation_type.
///
/// This should be called only after having made sure that all
/// preconditions are satisfied.
fn apply_mutations(&mut self, mutations: Vec<RegistryMutation>) {
    if mutations.is_empty() {
        // We should not increment the version if there is no
        // mutation, so that we keep the invariant that the
        // global version is the max of all versions in the store.
        return;
    }
    self.increment_version();
    self.apply_mutations_as_version(mutations, self.version);
}

//
https://github.com/dfinity/ic/blob/8ffffeb4be1/rs/registry/canister/src/registry.rs#L171-L185
/// Increments the latest version of the registry.
fn increment_version(&mut self) -> Version {
    self.version += 1;
    self.version
}
```

```

}

fn apply_mutations_as_version(
    &mut self,
    mut mutations: Vec<RegistryMutation>,
    version: Version,
) {
    // We sort entries by key to eliminate the difference between
    changelog
    // produced by the new version of the registry canister starting from
    v1
    // and the changelog recovered from the stable representation of the
    // original version that didn't support certification.
    mutations.sort_by(|l, r| l.key.cmp(&r.key));
    for m in mutations.iter_mut() {
        // We normalize all the INSERT/UPDATE/UPSERT operations to be just
        // UPSERTs. This serves 2 purposes:
        //
        // 1. This significantly simplifies reconstruction of the
    changelog
        // when we deserialize the registry from the original stable
        // representation.
        //
        // 2. This will play nicely with garbage collection: if an old
        // INSERT entry is removed, the newly connected clients won't
        // fail because of an UPDATE in the first survived entry with
    the
        // same key.
        m.mutation_type = match Type::from_i32(m.mutation_type).unwrap() {
            Type::Insert | Type::Update | Type::Upsert => Type::Upsert,
            Type::Delete => Type::Delete,
        } as i32;
    }

    let req = RegistryAtomicMutateRequest {
        mutations,
        preconditions: vec![],
    };
    let bytes = pb_encode(&req);

    // version 是之前已经 +1 的版本, 并转换成 [u8; 8]
    self.changelog.insert(version.into(), bytes);

    for mutation in req.mutations {
        (*self.store.entry(mutation.key).or_default()).push_back(RegistryValue {
            version,
            value: mutation.value,
            deletion_marker: mutation.mutation_type == Type::Delete as
i32,
        });
    }
}

```

```

pub enum Type {
    /// Key is expected to not exist in the registry at the current
    version.
    /// (This includes the case of a key that has existed in the past and
    /// later got deleted).
    /// The mutation will fail otherwise.
    Insert = 0,
    /// Key is expected to exist in the registry at the current version.
    /// The mutation will fail otherwise.
    Update = 1,
    /// Key is expected to exist in the registry at the current version.
    /// The mutation will fail otherwise.
    Delete = 2,
    /// If the key does not exist at the current version, it will be
    created.
    /// Otherwise, the value will be updated. The name is common in the
    /// database world, and means Update or Insert.
    Upsert = 4,
}

fn pb_encode(msg: &impl prost::Message) -> Vec<u8> {
    let mut buf = vec![];
    msg.encode(&mut buf).unwrap();
    buf
}

```

先对输入 mutations 进行了修改，然后对整个 RegistryAtomicMutateRequest 编码，并写入 changelog，然后把 RegistryMutation 按照 key 和 value 写入到 store。

总的来说是把参数里面的数据结构写入到 Registry，并且增加了 n 个版本号，插入了 n 条 changelog。在 store 里面增加了 n \* m 个数据对。其中 n 是最外面 mutations 的数据长度，m 是里面的 mutations 的平均长度。

需要注意的是

1. changelog 里面，将整个最里面的 mutations 编码作为 value，以 version 作为 key，插入到 changelog 的红黑树里面，已经把版本和对应的数据都保存好了，store 里面存的可能是为了方便查询的冗余数据，因为使用了 Map 结构。
2. RegistryAtomicMutateRequest 里面的 preconditions 字段没有使用，代码里面也是通过构造一个空的数组来填充的：`preconditions: vec![]`

最后，又调用了 `recertify_registry()`，根据刚刚更新的 version 和 changelog 得到 root\_hash，然后存到 `certified_data` 里面。

## canister\_pre\_upgrade

升级前函数，将系统的状态先提前 stable 保存一下，然后在升级的时候 stable 保存的数据是不会丢失的，然后升级之后有一个升级后函数讲数据恢复出来。

```
#[export_name = "canister_pre_upgrade"]
fn canister_pre_upgrade() {
    println!("{}", "canister_pre_upgrade", LOG_PREFIX);
    let registry = registry_mut();
    let mut serialized = Vec::new();
    let ss = RegistryCanisterStableStorage {
        registry: Some(registry.serializable_form()),
    };
    ss.encode(&mut serialized)
        .expect("Error serializing to stable.");
    stable::set(&serialized);
}
```

这个函数先打印一些信息出来，然后获取 Registry 示例 registry，其实这里没必要用 mut，因为并不会修改状态（删掉 mut 可以编译通过）。然后根据 registry 的信息构造了一个 RegistryCanisterStableStorage 数据，并且把这个数据编码到一个 buffer serialized 里面，然后调用 stable 的 set 接口保存这个数据。

来看一下 stable 的 set 是怎么存的

```
/// The wasm page size is 64KiB
const PAGE_SIZE: f64 = 64.0 * 1024.0;

/// We store all the data prepended by the length of data in the first 4
bytes
/// 0      4      4 + length
/// +-----+-----+----->
/// | length |   content   |   junk   |
/// +-----+-----+----->

const LENGTH_BYTES: u32 = 4;

//
https://github.com/dfinity/ic/blob/8fffeb4be1/rs/rust\_canisters/dfn\_core/s
rc/stable.rs#L37-L45
/// Sets the contents of the stable memory
pub fn set(content: &[u8]) {
    let len: u32 = content.len() as u32;
    ensure_capacity(len);
    unsafe {
        ic0::stable_write(LENGTH_BYTES, content.as_ptr() as u32, len);
    }
    set_length(len);
}

fn ensure_capacity(capacity_bytes: u32) {
    let required_pages = (f64::from(capacity_bytes + LENGTH_BYTES) /
PAGE_SIZE).ceil() as u32;
    let current_pages = unsafe { ic0::stable_size() };

    if required_pages > current_pages {
```

```

        let difference = required_pages - current_pages;
        unsafe {
            ic0::stable_grow(difference);
        };
    }
}

pub fn set_length(len: u32) {
    ensure_capacity(LENGTH_BYTES);
    let len_bytes = len.to_le_bytes();
    unsafe { ic0::stable_write(0, len_bytes.as_ptr() as u32, LENGTH_BYTES) }
}
}

```

有一个 wasm page 的概念，每个 page 64 KiB，ensure\_capacity 函数将 page 扩展到足够的空间，然后调用系统接口 stable\_write 从第 4 bytes 开始往里面写数据，最后往里面的前 4 个 bytes 存放数据的长度。

接口看一下 RegistryCanisterStableStorage 数据结构：

<b>RegistryCanisterStableStorage</b>			
<i>pub registry: Option&lt;</i>	<i>RegistryStableStorage</i>		
	<i>pub version: i32</i>		
	<i>pub deltas: Vec&lt;</i>	<i>RegistryDelta</i>	
		<i>pub key: Vec&lt;u8&gt;</i>	
		<i>RegistryValue</i>	
	<i>pub values: Vec&lt;</i>	<i>pub value: Vec&lt;u8&gt;</i>	
		<i>pub version: u64</i>	
		<i>pub deletion_marker: bool</i>	
	<i>pub changelog: Vec&lt;</i>	<i>ChangelogEntry</i>	
		<i>pub version: u64</i>	
		<i>pub encoded_mutation: Vec&lt;u8&gt;</i>	

因为是要把 Registry 结构存下来，所以和 Registry 结构比较类似。

```

/// Serializes the registry contents using the specified version of stable
/// representation.
fn serializable_form_at(&self, repr_version: ReprVersion) ->
RegistryStableStorage {
    match repr_version {
        ReprVersion::Version1 => RegistryStableStorage {
            version: repr_version as i32,
            deltas: vec![],
            changelog: self
                .changelog
                .iter()
                .map(|(encoded_version, bytes)| ChangelogEntry {
                    version: encoded_version.as_version(),
                    encoded_mutation: bytes.clone(),
                })
                .collect(),
        },
    },
}

```

```

    ReprVersion::Unspecified => RegistryStableStorage {
        version: repr_version as i32,
        deltas: self
            .store
            .iter()
            .map(|(key, values)| RegistryDelta {
                key: key.clone(),
                values: values.iter().cloned().collect(),
            })
            .collect(),
        changelog: vec![],
    },
}

pub fn serializable_form(&self) -> RegistryStableStorage {
    self.serializable_form_at(ReprVersion::Version1)
}

```

升级的时候，RegistryStableStorage 的 deltas 字段里面并没有存数据，而且 version 也是指定的，因此原来状态变量 Registry 里面的两个字段，version，store 都没有保存。不过，之前得出，changelog 里面反正有所有的信息：有 n 个树的节点，当前 version 就是 n，store 里面的值也可以通过遍历 RbTree 的节点来得到。

### canister\_post\_upgrade

升级后函数，同样的，执行 hook()，设置好标准输出以及 panic 时打印错误信息，发生 panic 所在的代码位置。

```

#[export_name = "canister_post_upgrade"]
fn canister_post_upgrade() {
    dfn_core::printer::hook();
    println!("{}", canister_post_upgrade, LOG_PREFIX);
    // Purposefully fail the upgrade if we can't find authz information.
    // Best to have a broken canister, which we can reinstall, than a
    // canister without authz information.
    let ss =
RegistryCanisterStableStorage::decode(stable::get().as_slice())
        .expect("Error decoding from stable.");
    let registry = registry_mut();
    registry.from_serializable_form(ss.registry.expect("Error decoding
from stable"));

    registry.check_global_invariants(&[]);
    recertify_registry();
}

```

在打印一些升级消息之后，接着从 stable 存储里面读到对应的数据，有一个 stable\_read 的系统接口，可以读到。先读前 4 字节，获取整个的长度，然后再根据长度获取存储的数据。

```

/// Gets the contents of the stable memory
pub fn get() -> Vec<u8> {
    let len = length();
    let mut out: Vec<u8> = vec![0; len as usize];
    unsafe {
        ic0::stable_read(out.as_mut_ptr() as u32, LENGTH_BYTES, len as
u32);
    }
    out
}

pub fn length() -> u32 {
    let mut len_bytes: [u8; 4] = [0; 4];
    unsafe {
        ic0::stable_read(len_bytes.as_mut_ptr() as u32, 0, LENGTH_BYTES);
    }
    u32::from_le_bytes(len_bytes)
}

```

看一下怎么从 RegistryStableStorage 恢复出 Registry 的。先检查 Registry 是不是所有的字段是不是空，或者为 0，然后得到 version。之前 version 是通过指定 ReprVersion::Version1，也就是执行 Version1 这个代码分支。先遍历 changelog 字段，获取到 entry 这个 ChangelogEntry 类型的实例，对 Registry 先填充无用信息，目的是把版本号对应起来，然后把 entry 里面的 encoded\_mutation 解码成 RegistryAtomicMutateRequest 类型，并将获取到的 mutations 写入 Registry。

而且从前面可以看到，entry（ChangelogEntry 类型）的 encoded\_mutation 是直接复制的 Registry 的节点（Node<K, V>）的 V，而这个 V 是直接来自 RegistryAtomicMutateRequest encode 过来的。所以 canister\_init 时所有的数据（包括 store 和 changelog）都会恢复过来。

```

/// Sets the content of the registry from its serialized representation.
///
/// Panics if not currently empty: this is only meant to be used in
/// canister_post_upgrade.
///
/// In post_upgrade, one should do as much verification as possible, and
/// panic for anything unexpected. Indeed, panicking here keeps the
/// pre-upgrade state unchanged, and gives the developer an opportunity
/// to try upgrading to a different wasm binary. As a corollary, any
/// lossy way of handling unexpected data must be banned in
/// post_upgrade.
pub fn from_serializable_form(&mut self, stable_repr:
RegistryStableStorage) {
    assert!(self.store.is_empty());
    assert!(self.changelog.is_empty());
    assert_eq!(self.version, 0);

    let repr_version =
ReprVersion::from_i32(stable_repr.version).unwrap_or_else(|| {
        panic!(
            "Version {} of stable registry representation is not supported

```



```

    by this canister",
        stable_repr.version
    )
});

match repr_version {
    ReprVersion::Version1 => {
        let mut current_version = 0;
        for entry in stable_repr.changelog {
            // Code to fix ICSUP-2589.
            // This fills in missing versions with empty entries so
            // that clients see an
            // unbroken sequence.
            // If the current version is different from the previous
            // version + 1, we
            // need to add empty records to fill out the missing
            // versions, to keep
            // the invariants that are present in the
            // client side.
            for i in current_version + 1..entry.version {
                let mutations = vec![RegistryMutation {
                    mutation_type: Type::Upsert as i32,
                    key: "_".into(),
                    value: "".into(),
                }];
                self.apply_mutations_as_version(mutations, i);
                self.version = i;
            }
            // End code to fix ICSUP-2589

            let req =
                RegistryAtomicMutateRequest::decode(&entry.encoded_mutation[..])
                    .unwrap_or_else(|err| {
                        panic!("Failed to decode mutation@{}: {}",
                            entry.version, err)
                    });
            self.apply_mutations_as_version(req.mutations,
                entry.version);
            self.version = entry.version;
            current_version = self.version;
        }
    }
    ReprVersion::Unspecified => {
        ...
    }
}
}

```

总之执行升级后函数，里面会填充一些无用数据，但是之前的数据都恢复过来了。

update\_authz

```
#[export_name = "canister_update update_authz"]
fn update_authz() {
    check_caller_is_root();
    over(candid_one, |_: Vec<MethodAuthzChange>| {
        println!(
            "{}update_authz was called. \
            This does not do anything, since the registry canister no
longer has any \
            function whose access is controlled using this mechanism.
\
            TODO(NNS1-413): Remove this once we are sure that there
are no callers.",
            LOG_PREFIX,
        );
    })
}
```

根据打印的信息，执行这个函数应该不会修改任何状态。确实也是这样，它首先调用 `check_caller_is_root`，通过系统接口 `msg_caller_copy` 获取的调用者的 `PrincipalId`，然后检查该 `PrincipalId` 是否和 Root 的 Canister id 对得上。

```
pub fn check_caller_is_root() {
    if caller() != PrincipalId::from(ic_nns_constants::ROOT_CANISTER_ID) {
        panic!("Only the root canister is allowed to call this method.");
    }
}

/// Returns the caller of the current call.
pub fn caller() -> PrincipalId {
    let len: u32 = unsafe { ic0::msg_caller_size() };
    let mut bytes = vec![0; len as usize];
    unsafe {
        ic0::msg_caller_copy(bytes.as_mut_ptr() as u32, 0, len);
    }
    PrincipalId::try_from(bytes).unwrap()
}
```

然后接着调用了一个 `over` 函数，输入两个函数作为参数，具体功能还没理得很清楚 (Todo)，大概为了方便使用不同的编码，比如有的对外接口是基于 `candid`，有的是 `protobuf`，做了个封装。这种方式调用函数后续会有很多。

```
/// Over allows you to create canister endpoints easily
/// ```no_run
/// # use dfn_core::over;
/// #[export_name = "canister_query happy_birthday"]
/// fn hb() {
///     fn happy_birthday((name, age): (String, u16)) -> String {
///         format!("Happy Birthday {}", name)
///     }
/// }
```

```

/// }
/// over(dfn_json::json, happy_birthday)
/// }
/// ```
///
/// This function always call `reply` unless it traps. If you need a
function
/// that may `reject` the call, use one of the `_may_reject` variant.
pub fn over<In, Out, F, Witness>(_: Witness, f: F)
where
  In: FromWire + NewType,
  Out: IntoWire + NewType,
  F: FnOnce(In::Inner) -> Out::Inner,
  Witness: FnOnce(Out, In::Inner) -> (Out::Inner, In),
{
  over_bytes(|inp| {
    // TODO(RPL-266) Rejecting instead of trapping seems more
    // natural for deserialization errors. Debate in Jira.
    let outer = In::from_bytes(inp).expect("Deserialization Failed");
    let input = outer.into_inner();
    let res = f(input);
    let output = Out::from_inner(res);
    Out::into_bytes(output).expect("Serialization Failed")
  })
}

```

## current\_authz

```

#[export_name = "canister_query current_authz"]
fn current_authz() {
  over(candid, |_: ()| {
    println!(
      "{}current_authz was called. \
        This always returns the default value, since the registry
canister's state no \
        longer contains a CanisterAuthzInfo. \
        TODO(NNS1-413): Remove this once we are sure that there
are no callers.",
      LOG_PREFIX,
    );
    CanisterAuthzInfo::default()
  })
}

```

同样根据打印的信息，执行这个函数应该不会返回有用信息。只是返回 CanisterAuthzInfo 的初始化数据。

CanisterAuthzInfo	
	MethodAuthzInfo
pub methods_authz: Vec<	pub method_name: String
	pub principal_ids: Vec<Vec<u8>>

get\_changes\_since

get\_certified\_changes\_since

get\_value

get\_latest\_version

get\_certified\_latest\_version

atomic\_mutate

atomic\_mutate(), 这个函数只允许 Governance 和 Root 这两个 canister 来调用。然后，把调用参数反序列化，如果反序列化成功，就对 Registry 的可变实例执行 maybe\_apply\_mutation\_internal() 来进行修改。然后调用 recertify\_registry() 来更新 certified\_data。最后把返回结果序列化返回给调用者。

```
#[export_name = "canister_update atomic_mutate"]
fn atomic_mutate() {
    let caller = dfn_core::api::caller();
    //
    // - The governance canister is always allowed to mutate the registry
    // - The root canister is also allowed, so that IDs of new NNS
    canisters can be
    // recorded.
    assert!(
        caller == GOVERNANCE_CANISTER_ID.get() || caller ==
        ROOT_CANISTER_ID.get(),
        "{}Principal {} is not authorized to call 'atomic_mutate'.",
        LOG_PREFIX,
        caller
    );
    println!("{}", "call 'atomic_mutate' from {}", LOG_PREFIX, caller);

    let response_pb = match deserialize_atomic_mutate_request(arg_data())
    {
        Ok(request_pb) => {

            registry_mut().maybe_apply_mutation_internal(request_pb.mutations);
            RegistryAtomicMutateResponse {
                errors: vec![],
                version: registry().latest_version(),
            }
        }
        Err(error) => {
            println!(
                "{}Received a mutate call, but the request could not de
                serialized due to: {}",
                LOG_PREFIX, error
            );
            let mut response_pb = RegistryAtomicMutateResponse::default();
            let error_pb = RegistryError {
                code: Code::MalformedMessage as i32,
```

```

        reason: error.to_string(),
        ..Default::default()
    };
    response_pb.errors.push(error_pb);
    response_pb
}
};

recertify_registry();

let bytes =
    serialize_atomic_mutate_response(response_pb).expect("Error serializing
response");
    reply(&bytes)
}

```

看一下反序列化函数入参是怎么做的，参数需要是 `Vec<u8>` 的类型，它能够解码成 `RegistryAtomicMutateRequest` 类型就行。然后序列号返回结果是把返回结果编码成 `Vec<u8>`。

```

/// Deserializes the arguments for a request to the atomic_mutate()
function in
/// the registry canister, from protobuf.
pub fn deserialize_atomic_mutate_request(
    request: Vec<u8>,
) -> Result<pb::v1::RegistryAtomicMutateRequest, Error> {
    match pb::v1::RegistryAtomicMutateRequest::decode(&request[..]) {
        Ok(request) => Ok(request),
        Err(error) => Err(Error::MalformedMessage(error.to_string())),
    }
}

/// Serializes a response for a atomic_mutate() request to the registry
/// canister.
//
// This uses the PB structs directly as this function is meant to
// be used in the registry canister only and thus there is no problem with
// leaking the PB structs to the rest of the code base.
pub fn serialize_atomic_mutate_response(
    response: pb::v1::RegistryAtomicMutateResponse,
) -> Result<Vec<u8>, Error> {
    let mut buf = Vec::new();
    match response.encode(&mut buf) {
        Ok(_) => Ok(buf),
        Err(error) => Err(Error::MalformedMessage(error.to_string())),
    }
}

```

这里能够直接 `.encode().decode()` 就把这些工作完成了，是因为 `RegistryAtomicMutateRequest` 这个 struct 已经实现好 `Message` 这个 trait，里面包含 `encode`, `decode` 方法。

```

/// Message corresponding to a list of mutations to apply, atomically, to
the
/// registry canister. If any of the mutations fails, the whole operation
will fail.
#[derive(candid::CandidType, candid::Deserialize, Eq)]
#[derive(Clone, PartialEq, ::prost::Message)]
pub struct RegistryAtomicMutateRequest {
    /// The set of mutations to apply to the registry.
    #[prost(message, repeated, tag="1")]
    pub mutations: ::prost::alloc::vec::Vec<RegistryMutation>,
    /// Preconditions at the key level.
    #[prost(message, repeated, tag="5")]
    pub preconditions: ::prost::alloc::vec::Vec<Precondition>,
}

```

最后看一下 `reply()`，它调用系统接口 `msg_reply_data_append()` 和 `msg_reply()` 来返回数据。这样除了函数本身有一个返回值之外，还可以通过这种方式返回数据。

```

/// Replies with the given byte array.
/// Note, currently we do not support chunkwise assembling of the
response.
/// Warning if you use this with an endpoint it will cause a trap due to
the
/// message trying to return multiple responses
pub fn reply(payload: &[u8]) {
    unsafe {
        ic0::msg_reply_data_append(payload.as_ptr() as u32, payload.len()
as u32);
        ic0::msg_reply();
    }
}

```

## bless\_replica\_version

```

#[export_name = "canister_update bless_replica_version"]
fn bless_replica_version() {
    check_caller_is_governance_and_log("bless_replica_version");
    over(candid_one, |payload: BlessReplicaVersionPayload| {
        bless_replica_version_(payload)
    });
}

#[candid_method(update, rename = "bless_replica_version")]
fn bless_replica_version_(payload: BlessReplicaVersionPayload) {
    registry_mut().do_bleas_replica_version(payload);
    recertify_registry();
}

```

这个函数只允许 Governance Canister 来调用。

```
fn check_caller_is_governance_and_log(method_name: &str) {
    let caller = dfn_core::api::caller();
    println!("{}", call: {} from: {}", LOG_PREFIX, method_name, caller);
    assert_eq!(
        caller,
        GOVERNANCE_CANISTER_ID.into(),
        "{}Principal: {} is not authorized to call this method: {}",
        LOG_PREFIX,
        caller,
        method_name
    );
}
```

```
/// Adds a new replica version to the registry and blesses it, i.e., adds
/// the version's ID to the list of blessed replica versions.
///
/// This method is called by the governance canister, after a proposal
/// for blessing a new replica version has been accepted.
pub fn do_bless_replica_version(&mut self, payload:
BlessReplicaVersionPayload) {
    println!("{}", do_bless_replica_version: {:?}, LOG_PREFIX, payload);

    let version = self.latest_version();
    // Get the current list
    let blessed_key = make_blessed_replica_version_key();
    let before_append = match self.get(blessed_key.as_bytes(), version) {
        Some(old_blessed_replica_version) => {
            decode_registry_value:<BlessedReplicaVersions>(
                old_blessed_replica_version.value.clone(),
            )
            .blessed_version_ids
        }
        None => vec![],
    };

    let after_append = {
        let mut copy = before_append.clone();
        copy.push(payload.replica_version_id.clone());
        copy
    };
    println!(
        "{}Blessed version before: {:?} and after: {:?}",
        LOG_PREFIX, before_append, after_append
    );

    let mutations = vec![
        // Register the new version (that is, insert the new
```

```

ReplicaVersionRecord)
    RegistryMutation {
        mutation_type: registry_mutation::Type::Insert as i32,
        key: make_replica_version_key(&payload.replica_version_id)
            .as_bytes()
            .to_vec(),
        value: encode_or_panic(&ReplicaVersionRecord {
            release_package_url: payload.release_package_url.clone(),
            release_package_sha256_hex:
payload.release_package_sha256_hex,
        }),
    },
    // Bless the new version (that is, update the list of blessed
versions)
    RegistryMutation {
        mutation_type: registry_mutation::Type::Upsert as i32,
        key: blessed_key.as_bytes().to_vec(),
        value: encode_or_panic(&BlessedReplicaVersions {
            blessed_version_ids: after_append,
        }),
    },
];

// Check invariants before applying mutations
self.maybe_apply_mutation_internal(mutations);
}

```

 BlessReplicaVersionPayload

update\_subnet\_replica\_version

```

#[export_name = "canister_update update_subnet_replica_version"]
fn update_subnet_replica_version() {
    check_caller_is_governance_and_log("update_subnet_replica_version");
    over(candid_one, |payload: UpdateSubnetReplicaVersionPayload| {
        registry_mut().do_update_subnet_replica_version(payload);
        recertify_registry();
    });
}

```

### *UpdateSubnetPayload*

```

pub subnet_id: PrincipalId
pub replica_version_id: String

```

```

pub fn do_update_subnet_replica_version(&mut self, payload:
UpdateSubnetReplicaVersionPayload) {
    println!(

```



```

        "{}do_update_subnet_replica_version: {:?}"',
        LOG_PREFIX, payload
    );

    check_replica_version_is_blessed(self, &payload.replica_version_id);

    // Get the subnet record
    let subnet_key =
make_subnet_record_key(SubnetId::from(payload.subnet_id));
    let mutation = match self.get(subnet_key.as_bytes(),
self.latest_version()) {
        Some(RegistryValue {
            value: subnet_record_vec,
            version: _,
            deletion_marker: _,
        }) => {
            let mut subnet_record =
                decode_registry_value::<SubnetRecord>
(subnet_record_vec.clone());
            subnet_record.replica_version_id = payload.replica_version_id;
            RegistryMutation {
                mutation_type: registry_mutation::Type::Update as i32,
                key: subnet_key.as_bytes().to_vec(),
                value: encode_or_panic(&subnet_record),
            }
        }
        None => panic!("Error while fetching the subnet record"),
    };

    let mutations = vec![mutation];

    // Check invariants before applying mutations
    self.maybe_apply_mutation_internal(mutations)
}

```

## update\_icp\_xdr\_conversion\_rate

```

#[export_name = "canister_update update_icp_xdr_conversion_rate"]
fn update_icp_xdr_conversion_rate() {
    check_caller_is_governance_and_log("update_icp_xdr_conversion_rate");
    over(candid_one, |payload: UpdateIcpXdrConversionRatePayload| {
        registry_mut().do_update_icp_xdr_conversion_rate(payload);
        recertify_registry();
    });
}

```

### *UpdateIcpXdrConversionRatePayload*

*pub data\_source: String*

*pub timestamp\_seconds: u64*

*pub xdr\_permyriad\_per\_icp: u64*

```
pub fn do_update_icp_xdr_conversion_rate(
    &mut self,
    payload: UpdateIcpXdrConversionRatePayload,
) {
    println!(
        "{}do_update_icp_xdr_conversion_rate: {:?}" ,
        LOG_PREFIX, payload
    );

    // If there is no ICP/XDR conversion rate, we have to Insert new one
    let mutations = vec![upsert(
        make_icp_xdr_conversion_rate_record_key()
            .as_bytes()
            .to_vec(),
        encode_or_panic:::<IcpXdrConversionRateRecord>(&payload.into()),
    )];

    // Check invariants before applying mutations
    self.maybe_apply_mutation_internal(mutations);
}
```

add\_node

add\_node\_operator

create\_subnet

add\_nodes\_to\_subnet

delete\_subnet

recover\_subnet

remove\_nodes\_from\_subnet

remove\_nodes

remove\_node\_directly

update\_node\_operator\_config

update\_subnet

clear\_provisional\_whitelist

set\_firewall\_config

update\_node\_rewards\_table

add\_or\_remove\_data\_centers

update\_unassigned\_nodes\_config

get\_node\_providers\_monthly\_xdr\_rewards

## Reference

---

- <https://smartcontracts.org/docs/interface-spec/index.html#certification>
- <https://smartcontracts.org/docs/interface-spec/index.html#canister-signatures>